


A Top-Level View of Computer Function and Interconnection

Chapter 3

Based on:
William Stallings
Computer Organization and Architecture, 11th Global Edition

Program Concept

- Hardwired systems are inflexible
- General purpose hardware can do different tasks, given correct control signals
- Instead of re-wiring, supply a new set of control signals



What is a Program?

- A sequence of steps
- For each step, an arithmetic or logical operation is done
- For each operation, a different set of control signals is needed


Function of Control Unit

- For each operation, a unique code is provided
 - e.g. ADD, MOVE
- A hardware segment accepts the code and issues the control signals

Components


- The Control Unit and Arithmetic and Logic Unit constitute the Central Processing Unit
- Data and instructions need to get into the system and results need to get out
 - Input / Output
- Temporary storage of code and results is needed
 - Main memory

Computer Components - Top Level View



Instruction Cycle

- Fetch cycle
 - Program Counter (PC) holds address of next instruction to fetch
 - Processor fetches instruction from memory location pointed to by PC
 - Increment PC
 - Instruction loaded into Instruction Register (IR)
 - Processor interprets instruction and performs required actions



- Execute cycle:
 - Processor-memory
 - Data transfer between CPU and main memory
 - Processor I/O
 - Data transfer between CPU and I/O module
 - Data processing
 - Arithmetic or logical operation on data
 - Control
 - Alteration of sequence of operations
 - e.g. jump
 - Combination of above

Characteristics of a Hypothetical Machine



(a) Instruction format



(b) Integer format

Program counter (PC) = Address of instruction

Instruction register (IR) = Instruction being executed

Accumulator (AC) = Temporary storage

(c) Internal CPU registers


0001 = Load AC from memory

0010 = Store AC to memory


0101 = Add to AC from memory

(d) Partial list of opcodes

Example of Program Execution



Instruction Cycle - State Diagram



Multiple Memory Access

- The execute cycle may involve more than one memory access
- For example the PDP-11 includes the instruction ADD B, A
 - PDP-11 has multiple registers rather than a single accumulator
- The instruction cycle would look like this:
 - Fetch the instruction ADD B, A
 - Read the contents of memory location A into a register
 - Read the contents of memory location B into a different register
 - Add the two values
 - Write the result from the processor to memory location A
- This results in a more complicated execute cycle


Interrupts

- Mechanism by which other modules (e.g. I/O) may **interrupt** normal sequence of processing
 - May or may not resume processing after handling the interrupt
- Way to improve processing efficiency
 - External devices are slower than the processor
 - E.g. don't want the processor to wait for the printer
 - Printer can interrupt the processor when it needs attention
- Used to stop a process when an error has occurred
 - Error may or may not be recoverable


Classes of Interrupts

- Program
 - e.g. overflow, division by zero
- Timer
 - generated by internal processor timer
 - used in pre-emptive multi-tasking
- I/O
 - from I/O controller
- Hardware failure
 - e.g. power failure, memory parity error

Transfer of Control via Interrupts




Program Flow Control





Interrupt Cycle

- Added to instruction cycle
- Processor checks for interrupt
 - Indicated by an interrupt signal
- If no interrupt, fetch next instruction
- If interrupt pending:
 - Suspend execution of current program
 - Save context
 - Set PC to start address of interrupt handler routine
 - Process interrupt
 - Restore context and continue interrupted program



Instruction Cycle with Interrupts



Program Timing - Short I/O Wait




Program Timing - Long I/O Wait



(a) No interrupts

(c) Interrupts; long I/O wait

Instruction Cycle (with Interrupts) - State Diagram




Multiple Interrupts


- Disable interrupts
 - Processor will ignore further interrupts whilst processing one interrupt
 - Interrupts remain pending and are checked after first interrupt has been processed
 - Interrupts handled in sequence as they occur
- Define priorities
 - Low priority interrupts can be interrupted by higher priority interrupts
 - When higher priority interrupt has been processed, processor returns to previous interrupt

Multiple Interrupts


Sequential



Nested



Time Sequence of Multiple Interrupts



I/O Function


- An I/O module (e.g., a disk controller) can exchange data directly with the processor
- Processor can read data from or write data to an I/O module
 - Processor identifies a specific device that is controlled by a particular I/O module
 - I/O instructions rather than memory referencing instructions
- In some cases it is desirable to allow I/O exchanges to occur directly with memory
 - The processor grants to an I/O module the authority to read from or write to memory so that the I/O memory transfer can occur without tying up the processor
 - The I/O module issues read or write commands to memory relieving the processor of responsibility for the exchange
 - This is known as direct memory access (DMA)

Connecting

- All the units must be connected
- Different type of connection for different type of unit
 - Memory
 - Input/Output
 - CPU


Memory Connection

- Receives and sends data
- Receives addresses (of locations)
- Receives control signals
 - Read
 - Write
 - Timing



Input / Output Connection


- Similar to memory from computer's viewpoint
- Output
 - Receive data from computer
 - Send data to peripheral
- Input
 - Receive data from peripheral
 - Send data to computer



- Receive control signals from computer
- Send control signals to peripherals
 - e.g. spin disk
- Receive addresses from computer
 - e.g. port number to identify peripheral
- Send interrupt signals (control)

CPU Connection


- Reads instruction and data
- Writes out data (after processing)
- Sends control signals to other units
- Receives (and acts on) interrupts



What is a Bus?

- There are a number of possible interconnection systems
- Single and multiple **bus** structures are most common
- Examples
 - Control/Address/Data bus (PC)
 - Unibus (DEC-PDP)
- A communication pathway connecting two or more devices
- Usually **broadcast**
- Often grouped
 - A number of channels in one bus
 - e.g. 32 bit data bus is 32 separate single bit channels

Bus Interconnection Scheme



Data Bus

- Carries data
 - Recall that there is no difference between “data” and “instruction” at this level
- Width is a key determinant of performance
 - 8, 16, 32, 64 bit

Address Bus

- Identify the source or destination of data
- E.g. CPU needs to read an instruction (data) from a given location in memory
- Bus width determines maximum memory capacity of system
 - E.g. 8080 has 16 bit address bus giving 64k address space

Control Bus

- Control and timing information
 - Memory read/write signal
 - Interrupt request
 - Clock signals
- If one module wishes to send data to another, it must:
 - obtain the use of the bus, and
 - transfer data via the bus
- If one module wishes to request data from another, it must:
 - obtain the use of the bus, and
 - transfer a request to the other module over the appropriate control and address lines
 - It must then wait for that second module to send the data

Problem 3.3

Consider a hypothetical 32-bit microprocessor having 32-bit instructions composed of two fields: the first byte contains the opcode and the remainder the immediate operand or an operand address.

- a) What is the maximum directly addressable memory capacity (in bytes)?

- b) Discuss the impact on the system speed if the microprocessor bus has:
 - i. a 32-bit local address bus and a 16-bit local data bus, or
 - ii. a 16-bit local address bus and a 16-bit local data bus.

- c) How many bits are needed for the PC and the IR?

Problem 3.4

Consider a hypothetical microprocessor generating a 16-bit address and having a 16-bit data bus.

- a) What is the maximum memory address space that the processor can access directly if it is connected to a 16-bit memory?

- b) What is the maximum memory address space that the processor can access directly if it is connected to an 8-bit memory?


- c) What architectural features will allow this microprocessor to access a separate I/O space?

- d) If an input and an output instruction can specify an 8-bit I/O port number, how many 8-bit I/O ports can the microprocessor support? How many 16-bit I/O ports? Explain.


Single Bus Problems

- Lots of devices on one bus leads to:
 - **Propagation delays**
 - Long data paths mean that **coordination** of bus use can adversely affect performance
 - If aggregate data transfer approaches bus capacity
- Most systems use **multiple** buses to overcome these problems

Traditional Bus Architecture



High Performance Bus Architecture



Bus Types

- **Dedicated**
 - Separate data and address lines
- **Multiplexed**
 - Shared lines
 - Address valid or data valid control line
 - Advantage:
 - Fewer lines
 - Disadvantages:
 - More complex control


Bus Arbitration

- More than one module controlling the bus
 - E.g. CPU and DMA controller
- Only one module may control bus at one time
- Arbitration may be
 - **Centralized**
 - Single hardware device controlling bus access
 - Bus controller
 - Arbiter
 - May be part of CPU or separate
 - **Distributed**
 - Each module may claim the bus
 - Control logic on all modules


Timing

- Coordination of events on bus
- **Synchronous**
 - Events determined by clock signals
 - Control bus includes clock line
 - A single 1-0 is a bus cycle
 - All devices can read clock line
 - Usually sync on leading edge
 - Usually a single cycle for an event
- **Asynchronous**


Synchronous Timing Diagram



Asynchronous Timing - Read Diagram



Asynchronous Timing - Write Diagram



Problem 3.11*

For a synchronous read operation, the memory module must place the data on the bus sufficiently ahead of the falling edge of the Read signal to allow for signal settling.

Assume a microprocessor bus is clocked at 20 MHz and that the Read signal begins to fall in the middle of the second half of T3.

- Determine the length of the memory read instruction cycle.
- When, at the latest, should memory data be placed on the bus? Allow 10 ns for the settling of data lines.

* Some values differ from question in book


Point-to-Point Interconnect

- Principal reason for change from bus to **point-to-point** was the electrical constraints encountered with increasing the frequency of wide synchronous buses
- At higher and higher data rates, it becomes difficult to perform the synchronization and arbitration functions in a timely fashion
- With multicore chips (multiple processors and significant memory on a single chip), the use of a conventional shared bus on the same chip magnified the difficulties of increasing bus data rate and reducing bus latency to keep up with the processors
- *Compared to shared bus, point-to-point interconnect has lower latency, higher data rate, and better scalability*

Quick Path Interconnect (QPI)


- Introduced in 2008 by Intel
- Multiple direct connections
 - Direct pairwise connections to other components eliminating the need for arbitration found in shared transmission systems
- Layered protocol architecture
 - Rather than simple use of control signals found in shared bus arrangements
- Packetized data transfer
 - Data sent as a sequence of packets each of which contains control headers and error control codes

Multicore Configuration using QPI




QPI Layers

- Physical layer:
 - Actual wires carrying the signals, as well as circuitry and logic to support ancillary features required in the transmission and receipt of the 1s and 0s
 - Unit of transfer is 20 bits, called a Phit (physical unit)
- Link layer:
 - Responsible for reliable transmission and flow control
 - Unit of transfer is an 80-bit Flit (flow control unit)
- Routing layer:
 - Provides the framework for directing packets through the fabric
- Protocol layer:
 - High-level set of rules for exchanging packets of data between devices.
 - Comprised of an integral number of Flits



Physical Interface




The QPI port consists of 84 individual links grouped as:

- Each data path consists of a pair of wires that transmits data one bit at a time; each pair is referred to as a lane
- There are 20 data lanes in each direction (transmit and receive), plus a clock lane in each direction
- QPI is capable of transmitting 20 bits in parallel in each direction
- The 20-bit unit is referred to as a Phit
- The form of transmission on each lane is known as differential signaling, or balanced transmission – signals are transmitted as a current that travels down one conductor and returns on the other
- The binary value depends on the voltage difference
- Typically, one line has a positive voltage value and the other line has zero voltage, and one line is associated with binary 1 and one line is associated with binary 0

QPI Physical Layer Lanes

- The physical layer also manages the translation between 80-bit flits and 20-bit phits using multilane distribution.
- The flits can be considered as a bit stream that is distributed across the data lanes in a round-robin fashion.
- This enables QPI to achieve very high data rates by implementing the physical link between two ports as multiple parallel channels.



QPI Link Layer

- Performs two key functions: flow control and error control
 - Operate on the level of the flit (flow control unit)
 - Each flit consists of a 72- bit message payload and an 8-bit error control code (cyclic redundancy check - CRC)
- Flow control function
 - Needed to ensure that a sending QPI entity does not overwhelm a receiving QPI entity by sending data faster than the receiver can process the data and clear buffers for more incoming data
- Error control function
 - Detects and recovers from bit errors


QPI Routing and Protocol Layers

- Routing Layer
 - Used to determine the course that a packet will traverse across the available system interconnects
 - Defined by firmware and describes the possible paths that a packet can follow
- Protocol Layer
 - Packet is defined as the unit of transfer
 - One key function performed is a cache coherency protocol - main memory values held in multiple caches are consistent
 - A typical data packet payload is a block of data being sent to or from a cache

Peripheral Component Interconnect (PCI)


- A popular high bandwidth, processor independent bus that can function as a mezzanine or peripheral bus
- Delivers better system performance for high speed I/O subsystems
- PCI Special Interest Group (SIG)
 - Created to develop further and maintain the compatibility of the PCI specifications
- **PCI Express (PCIe)**
 - Point-to-point interconnect scheme intended to replace bus-based schemes such as PCI
 - Key requirement is high capacity to support the needs of higher data rate I/O devices, such as Gigabit Ethernet
 - Support time dependent data streams

Typical Configuration using PCIe




PCIe Layers

- Physical Layer – the physical wires that carry the data
 - Also includes the circuitry and logic to support the transmission and receipt of 1's and 0's
- Data Link Layer – responsible for reliable transmission and flow control
 - Data packets at the data link layer called DLLPs
- Transaction Layer – generates and consumes data packets and manages the flow between two components
 - Data packets at the transaction layer called TLPs



PCIe Physical Layer

- PCIe operates in a similar manner to QPI with the exception that lanes are bidirectional and the number of lanes on a PCIe port can be 1, 4, 6, 16, or 32
- Bits sent to lanes in a round-robin scheme
- At each physical lane data are buffered and processed 16 bytes (128 bits) at a time
- Each block is encoded into a 130 bit code word for transmission
- There is no common clock; the receiver looks for transitions in the data for synchronization
- The extra 2 bits ensure that in a long sequence of 1's there are at least some 0's to provide these transitions



PCIe multilane distribution

PCIe Layers

Data Link Layer

- Sends packets between two devices
- Adds error checking and address bits to each TLP to ensure they arrive at the correct device
- Packets fall into the following categories:
 - Flow control – regulate the rate at which TLPs and DLLPs can be transmitted
 - Power management – manage power platform budgeting
 - ACK – acknowledge the receipt of a valid data packet
 - NAK – acknowledge the receipt of an invalid data packet; the packet must be resent

Transaction Layer

- Receives read and write requests from the software above the TL and creates request packets for transmission to a destination via the link layer
- Most transactions use a split transaction technique
 - A request packet is sent out by a source PCIe device which then waits for a response
- TL messages and some write transactions are posted transactions (meaning that no response is expected)
- TL packet format supports 32-bit memory addressing and extended 64-bit memory addressing